Yuxuan Chen

Malden, MA | yuxuan.chen.1031@gmail.com | +1(223)269-8058 | LinkedIn | GitHub

Education

Northeastern University, Boston, MA Expected Dec 2026
M.S. in Software Engineering Systems

Dickinson College, Carlisle, PA May 2024
Dual Degree, B.S. in Computer Science & B.S. in Quantitative Economics

Skills

Languages: Python, Java, SQL
AI/ML: LLM Systems, RAG, NLP, Prompt Engineering
Tools: FastAPI, PostgreSQL, MySQL, AWS, GCP, vLLM, Linux

Experience

Research Assistant, Northeastern University – Boston, MA May 2026 – Present

- Designing and developing a platform to evaluate how repeated exposure influences recommendation behavior across large language models
- Building backend services with Python, FastAPI, PostgreSQL, and OpenAI APIs to automate experiment execution, data collection, and evaluation workflows
- Implementing large-scale exposure simulations and ranking analysis pipelines to measure shifts in model recommendations and citation behavior
- Developing automated content generation and publishing workflows using LLMs, Playwright, and web automation tools for Generative Engine Optimization (GEO) research

Software Engineer Intern, Software Velocity Corporation – Boston, MA Sep 2025 - Jan 2026

- Designed and implemented a unified **LLM evaluation pipeline** for 10+ models across cloud GPU environments (L4/T4/A100) and multiple APIs
- Built scalable benchmarking workflows for **latency, throughput, and cost analysis**, enabling structured comparison of LLM performance across deployment configurations
- Developed **experiment tracking system** with run tagging (model, config, GPU spec, prompt set), ensuring **reproducibility and auditability** of AI experiments
- Orchestrated **concurrent multi-model inference workloads**, identifying system bottlenecks and improving end-to-end pipeline efficiency by **25–40%**
- Automated **cloud-based model deployment** and **execution workflows**, reducing setup time and improving reliability of large-scale LLM testing

Research Assistant, Dickinson College – Carlisle, PA Sep 2023 - Jul 2024

- Built **NLP pipelines** for large-scale text processing (300K+ documents, 100M+ data points) using **Python**
- Designed **data workflows** for structured corpus generation supporting downstream **NLP analysis** and modeling
- Performed **linguistic analysis** on **high-dimensional text data** to extract patterns for **model-driven insights**
- Developed **reusable querying** and **visualization components** for iterative **data workflows**

Data Analytics Intern, China Construction Bank – China May - Jul 2023

- Automated daily transaction CSV cleaning via Python (rename/dedup/date parsing), added validation checks and reusable scripts, cutting prep time **75% (2h → 30m)** for recurring analyses
- Analyzed **50,000+** transactions in SQL, building customer segmentation models and summary tables
- Partnered with **BI** to translate marketing questions into **SQL insights**, compiling **stakeholder-ready reporting**